

2.5 BLACS Broadcast Communication Routines

A broadcast sends data possessed by one process to all processes within a scope. Broadcast, much like point to point communication has two complementary operations. The process that owns the data to be broadcast issues a broadcast/send. All processes within the same scope must then issue the complementary broadcast/receive. The BLACS define that both broadcast/send and broadcast/receive are globally-blocking. This has several important implications. The first is that scoped operations (broadcasts or combines) must be strictly ordered, i.e., all processes within a scope must agree on the order of calls to separate scoped operations. This constraint falls in line with that already in place for the computation of message IDs, and is present in point to point communication as well. A less obvious result is that scoped operations with SCOPE='ALL' must be ordered with respect to any other scoped operation. This means that if there are two broadcasts to be done, one along a column, and one involving the entire process grid, all processes within the process column issuing the column broadcast must agree on which broadcast will be performed first.

The topology parameter determines how the messages involved in a distributed operation are sent. There are two main classes of topologies within the BLACS:

- Pipelining topologies (ring-based),
- Non-pipelining topologies (tree-based).

The BLACS have a special default topology to provide for minimizing the single operation time, and this is selected by using the default topology TOP=' '.

A pipeline for increasing direction can be obtained by setting TOP='INCREASING RING' or TOP='I'; a pipeline for codes owing across the processors in the opposite way can be obtained by setting TOP='DECREASING RING' or TOP='D'.

The syntaxes of the Broadcast/Send routines are the following:

```
vGEBS2D(ICONTXT,SCOPE,TOP,M,N,A,LDA)
void Cdgebs2d(int,char*,char*,int,int,double*,int);
void Cigebs2d(int,char*,char*,int,int,int*,int);

vTRBS2D(ICONTXT,SCOPE,TOP,UPLO,DIAG,M,N,A,LDA)
void Cdtrbs2d(int,char*,char*,char*,char*,int,int,double*,int lda)
ICONTXT (input) INTEGER
    The BLACS context handle.
SCOPE (input) CHARACTER*1
    Scope of processes to participate in operation.
TOP (input) CHARACTER*1
    Network topology to be emulated during communication.
UPLO (input) CHARACTER*1
    Indicates whether the matrix is upper (UPLO='U') or lower (UPLO='L')
    trapezoidal.
DIAG (input) CHARACTER*1
    Indicates whether the diagonal of the matrix is unit diagonal (DIAG='U'), and
    thus need not be communicated, or otherwise (DIAG='N').
M (input) INTEGER
```

The number of matrix rows to be broadcast.
N (input) INTEGER
The number of matrix columns to be broadcast.
A (input) TYPE array (LDA,N)
A pointer to the beginning of the (sub)array to be broadcast.
LDA (input) INTEGER
The leading dimension of the matrix A, i.e., the distance between two successive elements in a matrix row.

The syntaxes of the Broadcast/Receive routines are the following:

```
vGEBR2D(ICONTXT, SCOPE, TOP, M, N, A, LDA, RSRC, CSRC)
void Cdgebr2d(int, char*, char*, int, int, double*, int, int, int);

vTRBR2D(ICONTXT, SCOPE, TOP, UPLO, DIAG, M, N, A, LDA, RSRC, CSRC)
Cdtrbr2d(int, char*, char*, char*, char*, int, int, double*, int, int, int)
```

ICONTXT (input) INTEGER
The BLACS context handle.
SCOPE (input) CHARACTER*1
Scope of processes to participate in operation.
TOP (input) CHARACTER*1
Network topology to be emulated during communication.
UPLO (input) CHARACTER*1
Indicates whether the matrix is upper (UPLO='U') or lower (UPLO='L') trapezoidal.
DIAG (input) CHARACTER*1
Indicates whether the diagonal of the matrix is unit diagonal (DIAG='U'), and thus need not be communicated, or otherwise (DIAG='N').
M (input) INTEGER
The number of matrix rows to be broadcast.
N (input) INTEGER
The number of matrix columns to be broadcast.
A (output) TYPE array (LDA,N)
A pointer to the beginning of the (sub)array to be received/broadcast.
LDA (input) INTEGER
The leading dimension of the matrix A, i.e., the distance between two successive elements in a matrix row.
RSRC (input) INTEGER
Process row coordinate of the source of the broadcast.
CSRC (input) INTEGER
Process column coordinate of the source of the broadcast.

2.5.1 BLACS Combines Routines

In a combine operation, each participating process contributes data which is combined with other processes' data to produce a result. This result can be left on a particular process (called the destination process), or on all participating processes. If the result is left on only one process, we refer to the operation as a leave-on-one combine, and if the result is given to all participating processes we reference it as a leave-on-all combine. The BLACS package three kinds of combines are supported. They are element-wise summation, element-wise absolute value maximization, and element-wise absolute value minimization of

general rectangular arrays. Element-wise indicates that each element of the input array will be combined with the corresponding element from all other processes' arrays to produce the result.

The syntaxes of the Combines routines are the following:

```
vGSUM2D(ICONTXT,SCOPE, TOP,M,N,A,LDA,RDEST,CDEST)
void Cdgsum2d(int,char*,char*,int,int,double*,int,int,int);

vGAMX2D(ICONTXT,SCOPE, TOP,M,N,A,LDA,RA,CA,RCFLAG,RDEST,CDEST)
void Cdgamx2d(int,char*,char*,int,int,double*,int,int*,int*,
int,int,int);

vGAMN2D(ICONTXT,SCOPE, TOP,M,N,A,LDA,RA,CA,RCFLAG,RDEST,CDEST)
void Cdgamn2d(int,char*,char*,int,int,double*,int,int*,int*,
int,int,int);
```

ICONTXT (input) INTEGER

The BLACS context handle.

SCOPE (input) CHARACTER*1

Scope of processes to participate in operation.

TOP (input) CHARACTER*1

Network topology to be emulated during communication.

M (input) INTEGER

The number of matrix rows to be combined.

N (input) INTEGER

The number of matrix columns to be combined.

A (input/output) TYPE array (LDA,N)

A pointer to the beginning of the (sub)array to be combined.

LDA (input) INTEGER

The leading dimension of the matrix A, i.e., the distance between two successive elements in a matrix row.

RA (output) INTEGER array (RCFLAG,N)

If RCFLAG=-1, this array will not be referenced, and need not exist. Otherwise it is an integer array (of size at least RCFLAGxN) indicating the row index of the process that provided the maximum/minimum. If the calling process is not selected to receive the result, this array will contain intermediate (useless) results.

CA (output) INTEGER array (RCFLAG,N)

If RCFLAG=-1, this array will not be referenced, and need not exist. Otherwise it is an integer array (of size at least RCFLAGxN) indicating the column index of the process that provided the maximum/minimum. If the calling process is not selected to receive the result, this array will contain intermediate (useless) results.

RCFLAG (input) INTEGER

If RCFLAG=-1, then the arrays RA and CA are not referenced and need not exist. Otherwise, RCFLAG indicates the leading dimension of these arrays, and so must be $\geq M$.

RDEST (input) INTEGER

The process row coordinate of the process who should receive the result. If RDEST or CDEST= -1, all processes within the indicated scope receive the answer.

CDEST (input) INTEGER

The process column coordinate of the process who should receive the result. If RDEST or CDEST= -1, all processes within the indicated scope receive the answer.

Example 2.5.1. This example illustrates the modalities of using the BLACS functions for parallel calculation of dot product of two vectors).

```

/*
This program does a bone-headed parallel double precision dot product of two vectors.
Arguments are input on process {0,0}, and output everywhere else.
*/
#include <iostream>
#include <iomanip>
#include <fstream>
#include <sstream>
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <cmath>
extern "C" {
// Cblacs and Blas declarations
    void Cblacs_get(int, int, int*);
    void Cblacs_gridinit(int*, const char*, int, int);
    void Cblacs_gridinfo(int,int*,int*,int*,int*);
    void Cdgebs2d(int,char*,char*,int,int,double*,int);
    void Cigebs2d(int,char*,char*,int,int,int*,int);
    void Cdgebr2d(int,char*,char*,int,int,int*,int,int,int);
    void Cdgebr2d(int,char*,char*,int,int,double*,int,int,int);
    void Cdgsum2d(int,char*,char*,int,int,double*,int,int,int);
    double ddot_(int*,double*,int*,double*,int*);
    void Cblacs_gridexit(int);
    void Cblacs_exit(int);
}
int main(int argc, char **argv)
{
    int CONTXT, N=100,one = 1;
    double *X= NULL, *Y = NULL;
    char scope='A',top='I';
    // Local Scalars ..
    int IAM, NPROCS, NPROW, NPCOL, MYPROW, MYPCOL, I, LN;
    double LDDOT;
    int procrows = 3, proccols = 1;
    // Executable Statements ..
    // Find out what grid has been set up, and pretend it's 1-D
    Cblacs_get(0, 0, &CONTXT);
    Cblacs_gridinit(&CONTXT, "Row-major", procrows, proccols);
    Cblacs_gridinfo(CONTXT,&NPROW,&NPCOL,&MYPROW,&MYPCOL);
    IAM = MYPROW*NPCOL + MYPCOL;
    NPROCS = NPROW * NPCOL;
    X = new double[N];
    Y = new double[N];
    // Do bone-headed thing, and just send entire X and Y to * everyone
    if ((MYPROW==0) & (MYPCOL==0))
    {
        for(int i=0;i<N;i++)
        {
            {
                X[i]=rand()/1000000000.0;
                Y[i]=rand()/1000000000.0;
            }
            Cigebs2d(CONTXT,&scope,&top,one,one,&N, one);// Distribution of N
            Cdgebs2d(CONTXT, &scope ,&top,N,1,X, N);// Distribution of vector X
            Cdgebs2d(CONTXT, &scope ,&top,N,1,Y, N);// Distribution of vector Y
        }
        else
        {
            Cigebr2d(CONTXT,&scope,&top,1,1,&N,1,0,0);
            Cdgebr2d(CONTXT,&scope,&top,N,1,X,N,0,0);
            Cdgebr2d(CONTXT,&scope,&top,N,1,Y,N,0,0);
        }
    }
}

```

```
// Find out the number of local rows to multiply (LN), and where in vectors to start (I)
    LN = N/NPROCS;
    I = IAM*LN;
// Last process does any extra rows
    if (IAM == NPROCS-1) LN = LN + N%NPROCS;
// Figure dot product of my piece of X and Y
printf("IAM=%d, LN=%d I=%d for (%d,%d) proces \n",IAM,LN,I,MYPROW,MYPCOL);
    LDDOT = ddot_(&LN,&X[I],&one, &Y[I], &one);
/*== Fiecare proces determina produsul scalar al X[I] si Y[I] de dimensiunea (LN-1)*one+1 adica produsul scalar al vectorilor X[I:LN-1]*one+1 si Y[I:LN-1]*one+1
==== */
printf("Local dot product for (%d,%d) proces is %f \n",MYPROW,MYPCOL,LDDOT);
// Add local dot products to get global dot product;
// give all procs the answer
    Cdgsum2d(CONTXT,&scope, &top, 1,1, &LDDOT, 1, -1, 0);
if ((MYPROW==0) & (MYPCOL==0))
    printf("Global dot product for (%d,%d) proces is %f \n",MYPROW,MYPCOL,LDDOT);
Cblacs_gridexit(CONTXT);
Cblacs_exit(0);
}
```

Program results:

```
[MI_gr_TPS1@hpc]$ ./mpiCC_ScL Example2.5.1.cpp -o Example2.5.1.exe
[MI_gr_TPS1@hpc]$ /opt/openmpi/bin/mpirun -n 3 -host compute-0-10 Example2.5.1.exe
IAM=0, LN=33 I=0 for (0,0) process
IAM=1, LN=33 I=33 for (1,0) process
Local dot product for (1,0) process is 46.684624
IAM=2, LN=34 I=66 for (2,0) process
Local dot product for (2,0) process is 38.270497
Local dot product for (0,0) process is 40.759342
Global dot product for (0,0) process is 125.714463
```

The Blas function vDDOT computes a vector-vector dot product and have the syntaxis:

```
res = vdot(n, x, incx, y, incy)
res= vdot_(int*,double*,int*,double*,int*);
```

Input Parameters

n
INTEGER. Specifies the number of elements in vectors x and y.

x
*REAL for sdot, DOUBLE PRECISION for ddot
Array, DIMENSION at least (1+(n-1)*abs(incx)).*

incx
INTEGER. Specifies the increment for the elements of x.

y
*REAL for sdot DOUBLE PRECISION for ddot
Array, DIMENSION at least (1+(n-1)*abs(incy)).*

incy
INTEGER. Specifies the increment for the elements of y.

Output Parameters

res
*REAL for sdot, DOUBLE PRECISION for ddot
Contains the result of the dot product of x and y, if n is positive. Otherwise, res contains 0.*

The following example should demonstrate how the Cdgamx2d routine is used.

Example 2.5.2. *This example illustrates the modalities of using the BLACS functions for parallel carrying out the operations of reduction.*

Additional explanations. Let consider a series of matrices $\{A^k = \left\| a_{ij}^k \right\|_{i=1,m; j=1,n}^{k=1,p}$, where k represents the index of the process. Then the function Cdgamx2d(ctxt,&scope, &top, M,N,A_glob,M,RA,CA,RCFLAG,0,0) determinates the matrix $A_glob = \left\| a_{ij}^* \right\|_{i=1,m}^{j=1,n}$ where $a_{ij}^* = \max_{k=1,p} a_{ij}^k$.

```

/* =====
   This program illustrates the modalities of using the BLACS functions for parallel
   carrying out the operations of reduction.
   =====*/
#include <mpi.h>
#include <iostream>
#include <iomanip>
#include <string>
#include <fstream>
#include <sstream>
using namespace std;
extern "C" {
    /* Cblacs declarations */
    void Cblacs_pinfo(int*, int*);
    void Cblacs_get(int, int, int*);
    void Cblacs_gridinit(int*, const char*, int, int);
    void Cblacs_pcoord(int, int, int*, int*);
    void Cblacs_gridexit(int);
    void Cblacs_barrier(int, const char*);
    void Cdgamx2d(int, char*, char*, int, int, double*, int, int*, int*, int, int, int);
    void Cdgamn2d(int, char*, char*, int, int, double*, int, int*, int*, int, int, int);
}
int main(int argc, char **argv)
{
// Initalizearea
    int mpirank, namelen;
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    MPI_Init(&argc, &argv);
    MPI_Get_processor_name(processor_name, &namelen);
    MPI_Comm_rank(MPI_COMM_WORLD, &mpirank);
    bool mpiroot = (mpirank == 0);
    int iZERO = 0, RCFLAG;
    int M=4, N=4; // dimensiunea matricei (MxN)
    char scope='A', top='I';
    double *A_glob = NULL;
    int *RA, *CA;
    A_glob = new double[N*M];
    RA=new int[M*N];
    CA=new int[M*N];
// BLACS initialization (a network of 2x3 processes are created)
    int ctxt, myid, myrow, mycol, numproc;
    int procrows = 2, proccols = 3;
    Cblacs_pinfo(&myid, &numproc);
    Cblacs_get(0, 0, &ctxt);
    Cblacs_gridinit(&ctxt, "Row-major", procrows, proccols);
    Cblacs_pcoord(ctxt, myid, &myrow, &mycol);
// Toate procesele initializeaza matricea srand(mpirank) pentru ca generatorul de numere
// aleatoare sa fie diferit pentru fiecare proces
    for(int i=0; i<M*N; i++)
        A_glob[i]=mpirank; //+rand()/1000000000.0;
    if (mpirank ==0)

```

```

printf("====Rezultatul programului,%s \n",argv[0]);
MPI_Barrier(MPI_COMM_WORLD);
RCFLAG=M;
Cdgamx2d(ctxt,&scope, &top, M,N,A_glob,M,RA,CA,RCFLAG,0,0);
// Print matrix with maximum elements
if ((myrow==0)&(mycol==0))
{
cout << "=== Matrix with maximum elements:\n";
for (int r = 0; r < M; ++r) {
for (int c = 0; c < N; ++c) {
cout << setw(3) << *(A_glob + N*c + r) << " ";
}
cout << "\n";
}
for (int r = 0; r < M; ++r)
for (int c = 0; c < N; ++c)
{
// Tipar procesele care contin elementele maximele
printf("Elementul maximal A[%d,%d]=%f se contine in procesul (%d,%d)
\n",r,c,A_glob[r,c],RA[r,c],CA[r,c]);
}
}
// Se reinitializeaza matricea deoarece matricea cu elementele maximele a fost
// atribuita procesului (0,0)
for(int i=0;i<M*N;i++)
A_glob[i]=mpirank;//+rand()/1000000000.0;
Cdgamx2d(ctxt,&scope, &top,M,N,A_glob,M,RA,CA,RCFLAG,0,0);
// Se tipareste matricea cu emelentele minimele
if ((myrow==0)&(mycol==0))
{
cout << "===Matrice cu elementele miniale:\n";
for (int r = 0; r < M; ++r) {
for (int c = 0; c < N; ++c) {
cout << setw(3) << *(A_glob + N*c + r) << " ";
}
cout << "\n";
}
for (int r = 0; r < M; ++r)
for (int c = 0; c < N; ++c)
{
// Tipar procesele care contin elementele minimele
printf("Elementul manimal A[%d,%d]=%f se contine in procesul (%d,%d)
\n",r,c,A_glob[r,c],RA[r,c],CA[r,c]);
}
}
}
//Cleaning
/* Release resources */
delete[] A_glob;
Cblacs_gridexit(ctxt);
MPI_Finalize();
}

```

Rezultatele programului

```

[MI_gr_TPS1@hpc]$ ./mpiCC_ScL -o Example2.5.2.exe Example2.5.2.cpp
[MI_gr_TPS1@hpc]$/opt/openmpi/bin/mpirun -n 6 -host compute-0-0,compute-0-1
Example2.5.2.exe
===Matrice cu elementele maximele:
 5  5  5  5
 5  5  5  5
 5  5  5  5
 5  5  5  5
Elementul maximal A[0,0]=5.000000 se contine in procesul (1,2)
Elementul maximal A[0,1]=5.000000 se contine in procesul (1,2)
Elementul maximal A[0,2]=5.000000 se contine in procesul (1,2)
Elementul maximal A[0,3]=5.000000 se contine in procesul (1,2)

```

```

Elementul maximal A[1,0]=5.000000 se contine in procesul (1,2)
Elementul maximal A[1,1]=5.000000 se contine in procesul (1,2)
Elementul maximal A[1,2]=5.000000 se contine in procesul (1,2)
Elementul maximal A[1,3]=5.000000 se contine in procesul (1,2)
Elementul maximal A[2,0]=5.000000 se contine in procesul (1,2)
Elementul maximal A[2,1]=5.000000 se contine in procesul (1,2)
Elementul maximal A[2,2]=5.000000 se contine in procesul (1,2)
Elementul maximal A[2,3]=5.000000 se contine in procesul (1,2)
Elementul maximal A[3,0]=5.000000 se contine in procesul (1,2)
Elementul maximal A[3,1]=5.000000 se contine in procesul (1,2)
Elementul maximal A[3,2]=5.000000 se contine in procesul (1,2)
Elementul maximal A[3,3]=5.000000 se contine in procesul (1,2)
===Matrice cu elementele minimale:
  0  0  0  0
  0  0  0  0
  0  0  0  0
  0  0  0  0
Elementul manimal A[0,0]=0.000000 se contine in procesul (0,0)
Elementul manimal A[0,1]=0.000000 se contine in procesul (0,0)
Elementul manimal A[0,2]=0.000000 se contine in procesul (0,0)
Elementul manimal A[0,3]=0.000000 se contine in procesul (0,0)
Elementul manimal A[1,0]=0.000000 se contine in procesul (0,0)
Elementul manimal A[1,1]=0.000000 se contine in procesul (0,0)
Elementul manimal A[1,2]=0.000000 se contine in procesul (0,0)
Elementul manimal A[1,3]=0.000000 se contine in procesul (0,0)
Elementul manimal A[2,0]=0.000000 se contine in procesul (0,0)
Elementul manimal A[2,1]=0.000000 se contine in procesul (0,0)
Elementul manimal A[2,2]=0.000000 se contine in procesul (0,0)
Elementul manimal A[2,3]=0.000000 se contine in procesul (0,0)
Elementul manimal A[3,0]=0.000000 se contine in procesul (0,0)
Elementul manimal A[3,1]=0.000000 se contine in procesul (0,0)
Elementul manimal A[3,2]=0.000000 se contine in procesul (0,0)
Elementul manimal A[3,3]=0.000000 se contine in procesul (0,0)

```