

2.10 LU factorization of a general M-by-N distributed matrix.

In this section we first briefly describe the sequential block-partitioned versions of the dense LU, QR and Cholesky factorization routines of the LAPACK library. Since we also wish to discuss the parallel factorizations we describe the right-looking versions of the routines. The right-looking variants minimize data communication and distribute the computation across all processes. After describing the sequential factorizations the parallel versions will be discussed.

For the implementation of the parallel block partitioned algorithms in ScaLAPACK we assume that a matrix A is distributed over a $P \times Q$ process grid with a block cyclic distribution and a block size of $n_b \times n_b$ matching the block size of the algorithm. Thus each n_b -wide column (or row) panel lies in one column (row) of the process grid. In the LU, QR and Cholesky factorization routines in which the distribution of work becomes uneven as the computation progresses a larger block size results in greater load imbalance but reduces the frequency of communication between processes. There is therefore a tradeoff between load imbalance and communication startup cost which can be controlled by varying the block size. In addition to the load imbalance that arises as distributed data are eliminated from a computation load imbalance may also arise due to computational "hot spots" where certain processes have more work to do between synchronization points than others. This is the case for example in the LU factorization algorithm where partial pivoting is performed over rows in a single column of the process grid while the other processes are idle. Similarly the evaluation of each block row of the U matrix requires the solution of a lower triangular system across processes in a single row of the process grid. The effect of this type of load imbalance can be minimized through the choice of P and Q .

The LU factorization applies a sequence of Gaussian eliminations to form $A=PLU$, where A and L are $M \times N$ matrices, and U is an $N \times N$ matrix. L is unit lower

triangular (lower triangular with 1's on the main diagonal), U is upper triangular, and P is a permutation matrix, which is stored in a $\min(M;N)$ vector.

Constructing the Block LU Factorization

At the k -th step of the computation ($k=1,2,\dots$), it is assumed that the $m \times n$ submatrix of $A^{(k)}$ ($m=M-(k-1)n_b$, $n=N-(k-1)n_b$) is to be partitioned as follows

$$\begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} = P \begin{pmatrix} L_{11} & 0 \\ L_{21} & L_{22} \end{pmatrix} \begin{pmatrix} U_{11} & U_{12} \\ 0 & U_{22} \end{pmatrix} = P \begin{pmatrix} L_{11}U_{11} & L_{11}U_{12} \\ L_{21}U_{11} & L_{21}U_{12} + L_{22}U_{22} \end{pmatrix}$$

where the block A_{11} is $n_b \times (n-n_b)$, A_{12} is $(m-n_b) \times n_b$, and A_{22} is $(m-n_b) \times (n-n_b)$, L_{11} is a unit lower triangular matrix, and U_{11} is an upper triangular matrix.

At first a sequence of Gaussian eliminations is performed on the first $m \times n_b$ panel of $A^{(k)}$ (i. e. A_{11} and A_{21}). Once this is completed the matrices L_{11} , L_{21} and U_{11} are known and we can rearrange the block equations

$$\begin{aligned} U_{12} &\leftarrow (L_{11})^{-1} A_{12} \\ \tilde{A}_{22} &\leftarrow A_{22} - L_{21}U_{12} = L_{22}U_{22} \end{aligned}$$

The LU factorization can be done by recursively applying the steps outlined above to the $(m-n_b) \times (n-n_b)$ matrix \tilde{A}_{22} .

The computation of the above steps in the LAPACK routine **DGETRF** involves the following operations

1. **DGETF2**: Apply the LU factorization on an $m \times n_b$ column panel of A (i.e. A_{11} and A_{21})
 - [Repeat n_b times ($i = 1 \dots n_b$)]
 - **IDAMAX** find the (absolute) maximum element of the i -th column and its location
 - **DSWAP** interchange the i -th row with the row which holds the maximum
 - **DSCAL** scale the i -th column of the matrix
 - **DGER** update the trailing submatrix
2. **DLASWP** Apply row interchanges to the left and the right of the panel
3. **DTRSM** Compute the $n_b \times (n-n_b)$, row panel of

$$U, U_{12} \leftarrow (L_{11})^{-1} A_{12}$$

4. **DGEMM** Update the rest of the matrix A_{22} , $\tilde{A}_{22} \leftarrow A_{22} - L_{21}U_{12} = L_{22}U_{22}$

The corresponding parallel implementation of the ScaLAPACK routine **PDGETRF** proceeds as follows

1. **PDGETF2**: The current column of processes performs the LU factorization on an $m \times n_b$ column panel of A (i.e. A_{11} and A_{21})
 - [Repeat n_b times ($i = 1 \dots n_b$)]
 - **PDAMAX** find the (absolute) maximum value of the i -th column and its location (pivot information will be stored on the column of processes)
 - **PDLASWP** interchange the i -th row with the row which holds the maximum
 - **PDSCAL** scale the i -th column of the matrix
 - **PDGER** broadcast the i -th row columnwise ($(n_b - i)$ elements) in the current column of processes and update the trailing submatrix
 - Every process in the current process column broadcasts the same pivot information row wise to all columns of processes
2. **PDLASWP**: All processes apply row interchanges to the left and the right of the current panel
3. **PDTRSM**: L_{11} is broadcast along the current row of processes which converts the row panel A_{12} to U_{12}
4. **PDGEMM**: The column panel L_{21} is broadcast rowwise (*построчно*) across all columns of processes. The row panel U_{12} is broadcast columnwise down all rows of processes. Then all processes update their local portions of the matrix A_{22}

Syntaxes of the routine for use to compute the LU factorization of a general M-by-N distributed matrix

```
PvGETRF(M,N,A,IA,JA,DESCA,IPIV,INFO)
```

```
void pdgetrf_(int *m, int *n, double* a, int *ia, int *ja, int *desc_a, int  
             *ipiv, int *info);
```

Purpose: This routine forms the LU factorization of a general M-by-N distributed matrix $\text{sub}(A)=A(IA:IA+M-1,JA:JA+N-1)$ as $A=P*L*U$ where P is a permutation matrix, L is lower triangular with unit diagonal elements (lower trapezoidal if $M > N$) and U is upper triangular (upper trapezoidal if $M < N$). L and U are stored in $\text{sub}(A)$. The routine uses partial pivoting, with row interchanges.

Input Parameters

M (global) INTEGER.

The number of rows in the distributed submatrix $\text{sub}(A)$; $M \geq 0$.

N (global) INTEGER.

The number of columns in the distributed submatrix $\text{sub}(A)$; $N \geq 0$.

A (local) REAL for PSGETRF, DOUBLE PRECISION for PDGETRF, COMPLEX for PCGETRF, DOUBLE COMPLEX for PZGETRF.

Pointer into the local memory to an array of local dimension $(LLD_A, LOCC(JA+N-1))$. Contains the local pieces of the distributed matrix $\text{sub}(A)$ to be factored.

IA, JA (global) INTEGER.

The row and column indices in the global array A indicating the first row and the first column of the submatrix $A(IA:IA+N-1,JA:JA+N-1)$, respectively.

DESCA (global and local) INTEGER array, dimension $(dlen_)$.

The array descriptor for the distributed matrix A.

Output Parameters

A

Overwritten by local pieces of the factors L and U from the factorization $A = P*L*U$. The unit diagonal elements of L are not stored.

IPIV (local) INTEGER array.

The dimension of IPIV is $(LOC_R(M_A)+MB_A)$. This array contains the pivoting information: local row I was interchanged with global row $IPIV(I)$. This array is tied to the distributed matrix A .

INFO (global) INTEGER.

If $INFO=0$, the execution is successful.

$INFO < 0$: if the i -th argument is an array and the j -th entry had an illegal value, then $INFO = -(I*100+J)$; if the i -th argument is a scalar and had an illegal value, then $INFO = -I$.

If $INFO = I$, u_{ii} is 0. The factorization has been completed, but the factor U is exactly singular. Division by zero will occur if you use the factor U for solving a system of linear equations.

Example 2.10.1. *The using of function pdgetrf*

```
#include <stdio.h>
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include "mpi.h"
#include <iostream>
#include <iomanip>
#include <fstream>
#include <sstream>
using namespace std;
#define A(i,j) A[(i)*n+(j)]
#define L(i,j) L[(i)*n+(j)]
#define U(i,j) U[(i)*n+(j)]
#define A_distr(i,j) A_distr[(i)*k+(j)]
#define L_distr(i,j) L_distr[(i)*k+(j)]
#define U_distr(i,j) U_distr[(i)*k+(j)]
static int MAX( int a, int b ){
    if (a>b) return(a); else return(b);
}
extern "C"
{
void Cblacs_pinfo( int* mypnum, int* nprocs);
void Cblacs_get( int context, int request, int* value);
int Cblacs_gridinit( int* context, char * order, int np_row, int np_col);
void Cblacs_gridinfo( int context, int* np_row, int* np_col, int* my_row, int* my_col);
void Cblacs_gridexit( int context);
void Cblacs_barrier(int, const char*);
void Cblacs_exit( int error_code);
void Cblacs_pcoord(int, int, int*, int*);
int numroc_(int *n, int *nb, int *iproc, int *isrcproc, int *nprocs);
int indx12g_(int*, int*, int*, int*, int*);
```

```

void descinit_(int *desc, int *m, int *n, int *mb, int *nb, int *irsrc, int *icsrc, int *ictxt,
               int *lld, int *info);
void pdgeadd_(char *TRANS,int *M, int *N,double * ALPHA,double *A,int *IA,int *JA,int
               *DESCA,double *BETA,double *C, int *IC,int *JC,int *DESCC);
void pdgetrf_(int *m, int *n, double* a, int *ia, int *ja, int *desc_a, int *ipiv, int *info);
void pdgemm_( char *TRANSA,char *TRANSB,int *M,int *N,int *K,double *ALPHA,
               double *A,int *IA,int *JA,int *DESCA, double * B, int * IB, int * JB, int *
               DESCB,double * BETA,double * C, int * IC, int * JC, int * DESCC );
    } // extern "C"
int main(int argc, char **argv) {
int i_one = 1, i_zero = 0;
double zero=0.0E+0, one=1.0E+0;
int descA[9],descL[9],descU[9],descA_distr[9],descL_distr[9],descU_distr[9];
int iam,nprocs,nprow,npcol,myrow,mycol;
int m,n,mb,nb,mp,nq;
int i, j,mypnum;
int iloc,jloc;
int lld,lld_distr;
int ictxt,info,lwork;
int *ippiv;
double alpha, beta;
m=5; n=5;
mb=2; nb=2;
nprow=2; npcol=2;
double *A, *L,*U, *A_distr,*L_distr,*U_distr,*work,*tau;
Cblacs_pinfo(&iam,&nprocs);
Cblacs_get( -1, 0, &ictxt );
Cblacs_gridinit(&ictxt, "R", nprow, npcol );
Cblacs_gridinfo(ictxt, &nprow, &npcol, &myrow, &mycol );
if ( iam==0 ){
A = (double*)malloc(m*n*sizeof(double));
L = (double*)malloc(m*n*sizeof(double));
U = (double*)malloc(m*n*sizeof(double));
for(i=0;i<m;i++)
    for(j=0;j<n;j++)
        A[i*n+j]=(10*i+j);/(i+j);
}else{
A = NULL;
L = NULL;
U = NULL;
}
if (iam==0)
{
printf("===== RESULT OF THE PROGRAM %s \n",argv[0]);
cout << "Global matrix A:\n";
    for (i = 0; i < m; ++i) {
        for (j = 0; j < n; ++j) {
            cout << setw(5) << *(A + n*i + j) << " ";
        }
        cout << "\n";
    }
    cout << endl;
}

```

```

}
mp = numroc_( &m, &mb, &myrow, &i_zero, &nprow );
nq = numroc_( &n, &nb, &mycol, &i_zero, &npcol );
A_distr =(double*) malloc( mp*nq*sizeof(double));
L_distr =(double*) malloc( mp*nq*sizeof(double));
U_distr =(double*) malloc( mp*nq*sizeof(double));
ippiv = (int*) malloc( (mp+mb)*sizeof(int));
lld = MAX( numroc_( &n, &n, &myrow, &i_zero, &nprow ), 1 );
descinit_(descA, &m, &n, &m, &n, &i_zero, &i_zero, &ictxt, &lld, &info );
descinit_(descL, &m, &n, &m, &n, &i_zero, &i_zero, &ictxt, &lld, &info );
descinit_(descU, &m, &n, &m, &n, &i_zero, &i_zero, &ictxt, &lld, &info );
lld_distr = MAX( mp, 1 );
descinit_( descA_distr, &m, &n, &mb, &nb, &i_zero, &i_zero, &ictxt, &lld_distr, &info );
descinit_( descL_distr, &m, &n, &mb, &nb, &i_zero, &i_zero, &ictxt, &lld_distr, &info );
descinit_( descU_distr, &m, &n, &mb, &nb, &i_zero, &i_zero, &ictxt, &lld_distr, &info );
pdgeadd_( "N", &m, &n, &one, A, &i_one, &i_one, descA, &zero, A_distr, &i_one,
        &i_one, descA_distr );
pdgetrf_(&m, &n, A_distr, &i_one, &i_one, descA_distr, ippiv, &info);
pdgeadd_( "N", &m, &n, &one, A_distr, &i_one, &i_one, descA_distr, &zero, A, &i_one,
        &i_one, descA );
if (iam==0)
{
    cout << "Global matrix A (after use functionin pdgetrf_): \n";
    for (j = 0; j < n; ++j){
        //for (i = 0; i < m; ++i) {
        for (i = 0; i < m; ++i){
            // for (j = 0; j < n; ++j){
            cout << setw(5) << *(A + n*i + j) << " ";
        }
        cout << "\n";
    }
    cout << endl;
    for (j = 0; j < n; ++j){
        //for (i = 0; i < m; ++i){
        for (i = 0; i < m; ++i){
            //for (j = 0; j < n; ++j){
            if (i==j) L[i*n+j]=1;
            if (i<j) L[i*n+j]=A[i*n+j];
            if (i>j) L[i*n+j]=0;
            if (i>=j) U[i*n+j]=A[i*n+j]; else U[i*n+j]=0;
        }
    }
    cout << "Global matrix L :\n";
    for (j = 0; j < n; ++j)
    // for (i = 0; i < m; ++i)
    {
        for (i = 0; i < m; ++i){
            // for (j = 0; j < n; ++j){
            cout << setw(5) << *(L + n*i + j) << " ";
        }
        cout << "\n";
    }
}

```

```

        cout << endl;
cout << "Global matrix U :\n";
for (j = 0; j < n; ++j){
    for (i = 0; i < m; ++i){
        cout << setw(5) << *(U + n*i + j) << " ";
    }
    cout << "\n";
}
cout << endl;
}

if( myrow==0 && mycol==0 ){
free( A );free( L );free( U );
}
Cblacs_gridexit(ictxt );
Cblacs_exit( 0);
}

```

Program's results

```

[Hancu_B_S@hpc ScaLAPACK_Exemple_Curs_Online]$ ./mpiCC_ScL -o
Example2.10.1.exe Example2.10.1.cpp
[Hancu_B_S@hpc ScaLAPACK_Exemple_Curs_Online]$ /opt/openmpi/bin/mpirun -n 4 -
host compute-0-0 Example2.10.1.exe
===== RESULT OF THE PROGRAM Example2.10.1.exe
Global matrix A:
  0   1   2   3   4
10  11  12  13  14
20  21  22  23  24
30  31  32  33  34
40  41  42  43  44

Global matrix A (after use functioni pdgetrf_):
  4  14  24  34  44
  0  10  20  30  40
0.5  0.5   0   0   0
0.75 0.25   0   0   0
0.25 0.75   0   0   0

Global matrix L :
  1   0   0   0   0
  0   1   0   0   0
0.5  0.5   1   0   0
0.75 0.25   0   1   0
0.25 0.75   0   0   1

Global matrix U :
  4  14  24  34  44
  0  10  20  30  40
  0   0   0   0   0
  0   0   0   0   0
  0   0   0   0   0

```


[Hancu_B_S@hpc ScaLAPACK_Exemple_Curs_Online]\$

To verified the rezults we use the secvential varionat, i.e. the function dgetrf_.

```
#include <string>
#include <iostream>
#include <stdio.h>
#include <math.h>
#include <iomanip>
#include <fstream>
#include <sstream>
using namespace std;

extern "C" {
    void dgetrf_(int *m, int *n, double *a, int *lda, int *ipiv, int *info);
}

int main() {
    int i,j,m=5,n=5;
    int info;
    int LDA;
    double *A = new double[m*n];
    double *L = new double[m*n];
    double *U = new double[m*n];
    int *ipiv=new int[m];
    LDA=m;
    for(i=0;i<m;i++)
        for(j=0;j<n;j++)
            A[i*n+j]=(10*i+j); //i+j;
    cout << "Global matrix AA:\n";
    for (i = 0; i < m; ++i) {
        for (j = 0; j < n; ++j) {
            cout << setw(5) << *(A + n*i + j) << " ";
        }
        cout << "\n";
    }
    cout << endl;
    dgetrf_(&m, &n, A, &LDA, ipiv, &info);
    for (j = 0; j < n; ++j)
    {
        //for (j = 0; j < n; ++j)
        for (i = 0; i < m; ++i)
        {
            if (i==j) L[i*n+j]=1;
            if (i<j) L[i*n+j]=A[i*n+j];
            if (i>j) L[i*n+j]=0;
            if (i>=j) U[i*n+j]=A[i*n+j]; else U[i*n+j]=0;
        }
    }
    cout << "Global matrix L :\n";
    //for (i = 0; i < m; ++i)
    for (j = 0; j < n; ++j)
    {
```

```

        //for (j = 0; j < n; ++j)
        for (i = 0; i < m; ++i) {
            cout << setw(5) << *(L + n*i + j) << " ";
        }
        cout << "\n";
    }
    cout << endl;
cout << "Global matrix U :\n";
//for (i = 0; i < m; ++i)
    for (j = 0; j < n; ++j)
    {
        //for (j = 0; j < n; ++j)
        for (i = 0; i < m; ++i) {
            cout << setw(5) << *(U + n*i + j) << " ";
        }
        cout << "\n";
    }
    cout << endl;

cout << "LU decomposed matrix:\n " << endl;
//for(i=0 ; i<m ; i++)
    for(j=0 ; j<n ; j++){
        for(i=0 ; i<m ; i++)
//for(j=0 ; j<n ; j++)
cout << setw(5) << *(A + n*i + j) << " ";
        cout << endl;
    }
    cout << endl;
    delete[] A;
    return 0;
}

```

Rezultatele programului:

```

[[Hancu_B_S@hpc Pentru_Masterat]$ ./mpiCC_ScL -o dgetrf.exe dgetrf.cpp
[Hancu_B_S@hpc Pentru_Masterat]$ /opt/openmpi/bin/mpirun -n 1 -host compute-0-0,compute-0-4 dgetrf.exe

```

Global matrix AA:

```

0 1 2 3 4
10 11 12 13 14
20 21 22 23 24
30 31 32 33 34
40 41 42 43 44

```

Global matrix L :

```

1 0 0 0 0
0 1 0 0 0
0.5 0.5 1 0 0
0.75 0.25 0 1 0
0.25 0.75 0 0 1

```

Global matrix U :

```

4 14 24 34 44

```

```

0 10 20 30 40
0 0 0 0 0
0 0 0 0 0
0 0 0 0 0

```

LU decomposed matrix:

```

4 14 24 34 44
0 10 20 30 40
0.5 0.5 0 0 0
0.75 0.25 0 0 0
0.25 0.75 0 0 0

```

The auxiliar functions to LU factorization

SUBROUTINE PDLASWP(DIREC, ROWCOL, N, A, IA, JA, DESCA, K1, K2,IPIV)

.. Scalar Arguments ..

CHARACTER DIREC, ROWCOL

INTEGER IA, JA, K1, K2, N

.. Array Arguments ..

INTEGER DESCA(*), IPIV(*)

DOUBLE PRECISION A(*)

Purpose:

=====

PDLASWP performs a series of row or column interchanges on the distributed matrix $\text{sub}(A) = A(IA:IA+M-1, JA:JA+N-1)$. One interchange is initiated for each of rows or columns K1 through K2 of $\text{sub}(A)$. This routine assumes that the pivoting information has already been broadcast along the process row or column. Also note that this routine will only work for K1-K2 being in the same MB (or NB) block. If you want to pivot a full matrix, use PDLAPIV.

Arguments

=====

DIREC (global input) CHARACTER

Specifies in which order the permutation is applied:

= 'F' (Forward)

= 'B' (Backward)

ROWCOL (global input) CHARACTER

Specifies if the rows or columns are permuted:

= 'R' (Rows)

= 'C' (Columns)

N (global input) INTEGER

If ROWCOL = 'R', the length of the rows of the distributed matrix $A(*, JA:JA+N-1)$ to be permuted;

If ROWCOL = 'C', the length of the columns of the distributed matrix $A(IA:IA+N-1, *)$ to be permuted.

A (local input/local output) DOUBLE PRECISION pointer into the local memory to an array of dimension (LLD_A, *).

On entry, this array contains the local pieces of the banded matrix to which the row/columns interchanges will be applied. On exit the permuted distributed matrix.

- IA (global input) INTEGER
The row index in the global array A indicating the first row of sub(A).
- JA (global input) INTEGER
The column index in the global array A indicating the first column of sub(A).
- DESCA (global and local input) INTEGER array of dimension DLEN_.
The array descriptor for the distributed matrix A.
- K1 (global input) INTEGER
The first element of IPIV for which a row or column interchange will be done.
- K2 (global input) INTEGER
The last element of IPIV for which a row or column interchange will be done.
- IPIV (local input) INTEGER array, dimension LOCr(M_A)+MB_A for row pivoting and LOCc(N_A)+NB_A for column pivoting. This array is tied to the matrix A, IPIV(K) = L implies rows (or columns) K and L are to be interchanged.

void pdscal(int * N, double * ALPHA, double * X, int * IX, int * JX, int * DESCX, int * INCX)

.. Scalar Arguments ..

int * INCX, * IX, * JX, * N;
double * ALPHA;

.. Array Arguments ..

int * DESCX;
double * X;

endif

Purpose

=====

PDSCAL multiplies an n element subvector sub(X) by the alpha,
where

sub(X) denotes X(IX,JX:JX+N-1) if INCX = M_X,
X(IX:IX+N-1,JX) if INCX = 1 and INCX <> M_X.

Arguments

=====

N (global input) INTEGER

On entry, N specifies the length of the subvector sub(X)
N must be at least zero.

ALPHA (global input) DOUBLE PRECISION

On entry, ALPHA specifies the scalar alpha. When ALPHA
supplied as zero then the local entries of the array X
responding to the entries of the subvector sub(X) need
be set on input.

X (local input/local output) DOUBLE PRECISION array

On entry, X is an array of dimension (LLD_X, Kx), where

is at least $\text{MAX}(1, \text{Lr}(1, \text{IX}))$ when $\text{INCX} = \text{M_X}$
 $\text{MAX}(1, \text{Lr}(1, \text{IX}+\text{N}-1))$ otherwise, and, Kx is at
 $\text{Lc}(1, \text{JX}+\text{N}-1)$ when $\text{INCX} = \text{M_X}$ and $\text{Lc}(1, \text{JX})$ se.
 Before entry, this array contains the local entries of
 matrix X . On exit, $\text{sub}(X)$ is overwritten with the
 subvector.

IX (global input) INTEGER

On entry, IX specifies X's global row index, which points
 the beginning of the submatrix $\text{sub}(X)$.

JX (global input) INTEGER

On entry, JX specifies X's global column index, which
 to the beginning of the submatrix $\text{sub}(X)$.

DESCX (global and local input) INTEGER array

On entry, DESCX is an integer array of dimension DLEN_ .
 is the array descriptor for the matrix X .

INCX (global input) INTEGER

On entry, INCX specifies the global increment for
 elements of X . Only two values of INCX are supported
 this version, namely 1 and M_X . INCX must not be zero.

***void pdtrsm_(F_CHAR_T SIDE, F_CHAR_T UPLO, F_CHAR_T TRANS, F_CHAR_T
 DIAG, int M, int N, double ALPHA, double A, int IA, int JA, int
 DESCA, double B, int IB, int JB, int DESCB)***

.. Scalar Arguments ..

F_CHAR_T DIAG, SIDE, TRANS, UPLO;

int IA, IB, JA, JB, M, N;

double ALPHA;

.. Array Arguments ..

int DESCA, DESCB;

double A, B;

Purpose

=====

PDTRSM solves one of the matrix equations $\text{op}(\text{sub}(A))X = \alpha \text{sub}(B)$, or $X \text{op}(\text{sub}(A)) = \alpha \text{sub}(B)$,

where

$\text{sub}(A)$ denotes $A(\text{IA}:\text{IA}+\text{M}-1, \text{JA}:\text{JA}+\text{M}-1)$ if $\text{SIDE} = \text{'L'}$,

$A(\text{IA}:\text{IA}+\text{N}-1, \text{JA}:\text{JA}+\text{N}-1)$ if $\text{SIDE} = \text{'R'}$, and,

$\text{sub}(B)$ denotes $B(\text{IB}:\text{IB}+\text{M}-1, \text{JB}:\text{JB}+\text{N}-1)$.

Alpha is a scalar, X and $\text{sub}(B)$ are m by n submatrices, $\text{sub}(A)$
 a unit, or non-unit, upper or lower triangular submatrix and $\text{op}(Y)$
 is one of

$\text{op}(Y) = Y$ or $\text{op}(Y) = Y'$.

The submatrix X is overwritten on $\text{sub}(B)$.

Arguments

=====

SIDE (global input) CHARACTER1

On entry, SIDE specifies whether $\text{op}(\text{sub}(A))$ appears
 the left or right of X as follows:

SIDE = 'L' or 'l' $\text{op}(\text{sub}(A))X = \alpha \text{sub}(B)$,

SIDE = 'R' or 'r' $X \text{op}(\text{sub}(A)) = \alpha \text{sub}(B)$.

UPLO (global input) CHARACTER1

On entry, UPLO specifies whether the submatrix $\text{sub}(A)$

an upper or lower triangular submatrix as follows:

UPLO = 'U' or 'u' sub(A) is an upper triangular submatrix,

UPLO = 'L' or 'l' sub(A) is a lower triangular submatrix.

TRANSA (global input) CHARACTER1

On entry, TRANSA specifies the form of op(sub(A)) to be used in the matrix multiplication as follows:

TRANSA = 'N' or 'n' op(sub(A)) = sub(A),

TRANSA = 'T' or 't' op(sub(A)) = sub(A)',

TRANSA = 'C' or 'c' op(sub(A)) = sub(A)'.

DIAG (global input) CHARACTER1

On entry, DIAG specifies whether or not sub(A) is triangular as follows:

DIAG = 'U' or 'u' sub(A) is assumed to be unit gular,

DIAG = 'N' or 'n' sub(A) is not assumed to be unit angular.

M (global input) INTEGER

On entry, M specifies the number of rows of the sub(B). M must be at least zero.

N (global input) INTEGER

On entry, N specifies the number of columns of the sub(B). N must be at least zero.

ALPHA (global input) DOUBLE PRECISION

On entry, ALPHA specifies the scalar alpha. When ALPHA is supplied as zero then the local entries of the array corresponding to the entries of the submatrix sub(B) are not set on input.

A (local input) DOUBLE PRECISION array

On entry, A is an array of dimension (LLD_A, Ka), where Ka is at least Lc(1, JA+M-1) when SIDE = 'L' or 'l' and is at least Lc(1, JA+N-1) otherwise. Before entry, this array contains the local entries of the matrix A.

Before entry with UPLO = 'U' or 'u', this array contains local entries corresponding to the entries of the upper triangular submatrix sub(A), and the local entries corresponding to the entries of the strictly lower triangular part of the submatrix sub(A) are not referenced.

Before entry with UPLO = 'L' or 'l', this array contains local entries corresponding to the entries of the lower triangular submatrix sub(A), and the local entries corresponding to the entries of the strictly upper triangular part of the submatrix sub(A) are not referenced.

Note that when DIAG = 'U' or 'u', the local entries corresponding to the diagonal elements of the submatrix sub(A) are not referenced either, but are assumed to be unity.

IA (global input) INTEGER

On entry, IA specifies A's global row index, which points to the beginning of the submatrix sub(A).

JA (global input) INTEGER

On entry, JA specifies A's global column index, which

to the beginning of the submatrix sub(A).

DESCA (global and local input) INTEGER array

On entry, DESCA is an integer array of dimension DLEN_.
is the array descriptor for the matrix A.

B (local input/local output) DOUBLE PRECISION array

On entry, B is an array of dimension (LLD_B, Kb), where Kb
at least $Lc(1, JB+N-1)$. Before entry, this array ns
the local entries of the matrix B.

On exit, the local entries of this array corresponding to
to the entries of the submatrix sub(B) are overwritten
the local entries of the m by n solution submatrix.

IB (global input) INTEGER

On entry, IB specifies B's global row index, which points
the beginning of the submatrix sub(B).

JB (global input) INTEGER

On entry, JB specifies B's global column index, which
to the beginning of the submatrix sub(B).

DESCB (global and local input) INTEGER array

On entry, DESCB is an integer array of dimension DLEN_.
is the array descriptor for the matrix B.